



Lab 14.2 — Ansible Roles (Deep Dive & Best Practices)

Author: Sandeep Kumar Sharma



Learning Objectives

In this lab you will learn: - Full role anatomy and each directory purpose in detail - Variable precedence and where to put variables (`defaults` vs `vars`) - How handlers, files, templates, and tasks interact inside a role - Role parameters and how to pass variables to roles - Using `meta/main.yml` for role dependencies and metadata - Importing/including roles dynamically (`include_role`, `import_role`) - Best practices for authoring reusable roles - How to test a role quickly using a small playbook



Learning Outcomes

After completing this lab you will be able to: - Create production-ready roles with clear separation of concerns - Decide where to place variables and why - Use handlers and templates correctly inside roles - Compose roles with dependencies and reuse them across projects - Apply tags and limits to execute parts of a role



Role Anatomy — Detailed Explanation

A role is a filesystem layout which Ansible understands. A typical role looks like:

```
my_role/
├── defaults/
│   └── main.yml      # Lowest precedence variables (safe to override)
├── vars/
│   └── main.yml      # Higher precedence variables (harder to override)
├── tasks/
│   └── main.yml      # Task list executed by the role
├── handlers/
│   └── main.yml      # Handlers triggered by notify
├── templates/
│   └── main.j2        # Jinja2 templates (.j2)
├── files/
│   └── main.conf      # Static files to copy with the copy module
└── meta/
    └── main.yml      # Role metadata and dependencies
```

```
└── tests/
    └── test.yml          # Optional small playbook used for manual testing
    └── README.md          # Role documentation (highly recommended)
```



Purpose of Each Directory (Expanded)

- **defaults/** — Variables here have the *lowest* precedence. Ideal for configurable defaults that users can override via inventory, extra vars, or group_vars.
- **vars/** — Variables here have *higher* precedence than defaults and are generally used for internal values that should not be easily overridden.
- **tasks/** — The actions the role performs. Keep `main.yml` concise and break complex logic into multiple task files included with `include_tasks`.
- **handlers/** — Actions that run only when notified by tasks (e.g., restart service after config change).
- **templates/** — Dynamic configuration files using Jinja2. Use `template` module to render them.
- **files/** — Static files you want to copy as-is with the `copy` module.
- **meta/** — Contains `main.yml` with metadata (author, license) and `dependencies` — other roles that must be applied first.
- **tests/** — Small playbooks to exercise the role manually; recommended before adding automated tests.
- **README.md** — Document variables, examples, and usage instructions.

Variable Precedence — Quick Reference

From lowest to highest precedence (only relevant items shown): 1. `defaults` role variables (`defaults/main.yml`) 2. inventory `group_vars` 3. inventory `host_vars` 4. playbook `vars_files` 5. playbook `vars` 6. extra vars (`ansible-playbook -e`) — highest priority

Note: `vars/main.yml` inside a role has higher precedence than `defaults/main.yml` and will typically override defaults. Use `defaults` for values you expect users to change.



Role Variables — Where to put what

- Use **defaults/** for tunable settings (package name, ports, paths).
- Use **vars/** for internal constants or computed values that you do not want overridden easily.
- Do NOT place secrets in `defaults` or `vars` — use Ansible Vault or external secret managers.



Handlers, Templates & Files Interaction

- Tasks modify or deploy a file/template and `notify` a handler.
- Handlers are defined in `handlers/main.yml` inside the role.
- Templates live under `templates/` and are deployed with the `template` module.
- Files that must be copied untouched go to `files/` and are used with the `copy` module.

Example sequence inside a role: 1. Install package 2. Deploy template to config path (template) 3. Notify handler to restart service (notify) 4. Handler runs once at the end of the role if any notify triggered



Passing Variables to Roles (Role Parameters)

You can pass variables to a role from the playbook:

```
- hosts: dev
  roles:
    - role: my_role
      var1: value1
      var2: value2
```

These variables behave like play-level vars and will override `defaults` but can be overridden by extra vars.



Role Dependencies (`meta/main.yml`)

Define dependencies in `meta/main.yml` to ensure roles are applied in the correct order. Example:

```
# meta/main.yml
---
dependencies:
  - { role: common, some_var: 10 }
  - role: firewall
```

When the role with dependencies is applied, Ansible will run the dependent roles first.



Including & Importing Roles Dynamically

- `import_role` — static, evaluated during playbook parsing

- `include_role` — dynamic, evaluated during runtime and supports `when` conditions

Example `include_role` with condition:

```
- name: Include role only when needed
  include_role:
    name: backup_role
    when: backup_enabled | default(false)
```

Use `include_role` when you want runtime decision making.



Best Practices for Authoring Roles

1. **Keep roles small and single-purpose.** One role should do one job (e.g., `webserver`, `database`).
 2. **Name variables with the role prefix** to avoid clashes: `webserver_port`, `webserver_user`.
 3. **Provide sensible defaults** in `defaults/main.yml` and document them in `README.md`.
 4. **Avoid hard-coded values;** prefer variables and defaults.
 5. **Use handlers** for service restarts and ensure they are idempotent.
 6. **Keep tasks idempotent** — consider `creates`, `removes`, or `when` checks.
 7. **Use tags** in tasks to allow selective execution during development.
 8. **Document the role:** variables, example usage, supported platforms, author.
 9. **Store secrets outside the role** and use Vault when necessary.
 10. **Use tests/ playbooks** to validate role behavior quickly.
-



Example — Full Role: `webserver_role` (Files & Snippets)

`defaults/main.yml`

```
---
webserver_package: "{{ 'apache2' if ansible_facts['os_family'] == 'Debian' else
'httpd' }}"
webserver_service: "{{ webserver_package }}"
webserver_listen_port: 80
```

`tasks/main.yml`

```

---
- name: Install web package
  package:
    name: "{{ webserver_package }}"
    state: present

- name: Deploy web config
  template:
    src: web.conf.j2
    dest: "/etc/{{ 'apache2' if ansible_facts['os_family'] == 'Debian' else 'httpd' }}/conf.d/web.conf"
    notify: restart webserver

- name: Ensure document root exists
  file:
    path: /var/www/html
    state: directory
    mode: '0755'

- name: Deploy test index.html
  copy:
    src: index.html
    dest: /var/www/html/index.html

```

handlers/main.yml

```

---
- name: restart webserver
  service:
    name: "{{ webserver_service }}"
    state: restarted

```

templates/web.conf.j2

```

Listen {{ webserver_listen_port }}
ServerName {{ inventory_hostname }}

<VirtualHost *:{{ webserver_listen_port }}>
  DocumentRoot /var/www/html
</VirtualHost>

```

files/index.html

```
<html>
  <body>
    <h1>Deployed by Ansible Role</h1>
  </body>
</html>
```

meta/main.yml

```
---
dependencies:
  - role: common
```

tests/test.yml (simple test playbook)

```
---
- hosts: dev
  roles:
    - role: webserver_role
```



How to run the role quickly

From the repository root:

```
ansible-playbook tests/test.yml
```

To pass custom variables:

```
ansible-playbook tests/test.yml -e "webserver_listen_port=8080"
```



Role Versioning & Distribution

- Use semantic versioning in your role repository tags (v1.0.0).
- Publish reusable roles to Ansible Galaxy or an internal Galaxy server for team reuse.
- Add `meta/main.yml` fields like `galaxy_info` to improve discovery.



Testing Roles (Manual & Automated)

- Manual: use `tests/test.yml` as above.
 - Automated: use Molecule (optional) for scenario-based tests and linting.
-



Documentation — What to include in README.md

- Role purpose
 - Supported OS/families
 - Variables (defaults + description)
 - Example usage
 - Dependencies
 - Author and license
-



Lab Summary

This lab covered a deep, practical view of Ansible roles: structure, variable placement, handlers, templates, dependencies, best practices, and testing.

Author: Sandeep Kumar Sharma